
SEE-2-SOUND : Zero-Shot Spatial Environment-to-Spatial Sound

Rishit Dagli¹ Shivesh Prakash¹ Robert Wu¹ Housman Khosravani^{1,2,3}

¹University of Toronto, Toronto, Canada

²Temerty Centre for Artificial Intelligence Research and Education in Medicine

³Sunnybrook Research Institute, Toronto, Canada

rishit@cs.toronto.edu, shivesh.prakash@mail.utoronto.ca,

rupert@cs.toronto.edu, h.khosravani@utoronto.ca

see2sound.github.io

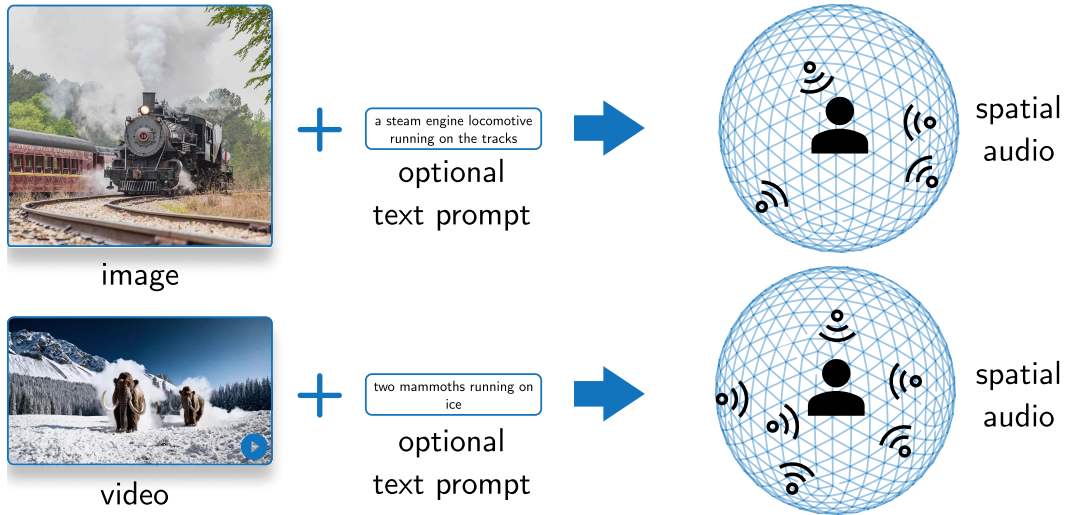


Figure 1: SEE-2-SOUND enables spatial audio creation from real or generated visual content. Our method takes in an image or a sequence of frames (video) and generates spatial audio with a variable number of different sources. The generated audio could also be conditioned on a text prompt.

Abstract

Generating combined visual and auditory sensory experiences is critical for the consumption of immersive content. Recent advances in neural generative models have enabled the creation of high-resolution content across multiple modalities such as images, text, speech, and videos. Despite these successes, there remains a significant gap in the generation of high-quality spatial audio that complements generated visual content. Furthermore, current audio generation models excel in either generating natural audio or speech or music but fall short in integrating spatial audio cues necessary for immersive experiences. In this work, we introduce SEE-2-SOUND, a zero-shot approach that decomposes the task into (1) identifying visual regions of interest; (2) locating these elements in 3D space; (3) generating mono-audio for each; and (4) integrating them into spatial audio. Using our

framework, we demonstrate compelling results for generating spatial audio for high-quality videos, images, and dynamic images from the internet, as well as media generated by learned approaches.

1 Introduction

Human visual and auditory systems inherently depend on synchronized spatial audio cues for an immersive experience and contextual understanding of video content. Modern generative models have been able to create compelling high-resolution digital content across various modalities. These models have been successful in generating images [61, 58], text [55, 1], speech [75, 57], audio [44, 35], 3D scenes [40, 53], 4D (dynamic 3D) scenes [68, 42], and videos [67, 13]. Although generating audio from text [30, 18] - and more recently, images [72] - has become a reality, generating spatial audio is still a challenge. Unlike these earlier works, we are interested in solving the specific problem of generating high-quality spatial audio or surround sound to accompany images, animated images (GIFs), and videos. Such a model could be particularly useful to pair with any kind of video generation or image generation techniques as well as on real images and videos.

Audio generation is a long-standing challenge, often subdivided into generating speech, music, and sound effects. Most audio generation methods include inductive biases for each of these tasks, often due to the diversity between the generation of speech, music, and other types of audio [11, 2, 43, 70]. There have also been newer approaches that develop unified representations [44] to generate complex audio as well. However, none of these currently emphasize the generation of rich spatial audio. The few related works that perform image-to-audio [72, 71] are closest to our work. However, these also face the same problem of not being able to generate rich spatial audio. Currently, there also exist a few learned approaches that generate spatial audio [47, 41] but these are often based on having access to the captured audio; following which these approaches localized the audio in a viewing sphere to produce spatial audio for a captured video. However, none of these approaches aim to generate accompanying spatial audio for images, animated images (GIFs), and videos, which remains an open problem.

Towards this goal, we propose SEE-2-SOUND, a training-free method that approaches this specific problem by employing the foundational models, CoDi [72], DepthAnything [84], and Segment Anything [33]. The pipeline works by explicitly factorizing this generative problem into separate problems: (1) identifying regions of interest in the input; (2) finding the 3D location on the viewing sphere for a visual element; (3) generating audio from the regions of interest, and (4) simulating these multiple mono-audio outputs to generate spatial audio. We then explicitly compose these component models to generate spatial audio conditioned on the visual input (e.g. images, animated images (GIFs), or videos) as shown in Figure 1. As our method is compositional and zero-shot, we hope to see SEE-2-SOUND widely used or adapted for content creation and enhancement as well as accessibility.

By generating spatial audio that accurately corresponds to visual elements, we can significantly enhance the immersion of images and videos, creating more realistic and engaging user experiences with digital content. SEE-2-SOUND presents a step towards true complete video and image generation, by also incorporating non-verbal audio that accompanies visual generation models [67, 13, 8, 9, 46, 79]. Additionally, our pipeline can also be employed on real visual content.

Our contributions can be summarized as follows:

- We present the challenging new task of generating spatial audio from images or videos. To the best of our knowledge, our work is the first of its kind that tackles this challenging problem.
- We propose SEE-2-SOUND, a novel compositional pipeline that builds on top of and uses off-the-shelf models for segmentation, depth estimation, and audio generation. Due to the nature of our approach, our method is zero-shot in nature and thus is very feasible to use.
- We evaluate SEE-2-SOUND both quantitatively and qualitatively through similarity comparisons and a user study, respectively. To evaluate this method quantitatively we also propose a novel evaluation method for this task.

We have open-sourced our code and models to promote the accessibility and reproducibility of our work. We also hope the extensibility of our work allows users to easily modify or replace individual modules within our approach. Our assets can be found at: <https://github.com/see2sound/see2sound>.

2 Related works

The audio generation field has been gaining much attention lately, primarily in models that create audio conditioned on different modalities. From speech synthesis to music generation and sound effects creation, past efforts have explored numerous aspects of auditory information synthesis. Learning algorithms have led to ever-improving techniques in the models for generating high-quality, diverse audio. This surge in interest has fueled research efforts that push the boundaries of audio generation.

2.1 Conditional Audio Generation

Recent work in conditional audio generation has made significant progress using latent diffusion models [43] such as AudioLDM (2) [43, 44] and Make-An-Audio [30]. Audiogen [35] was one of the first to create a text-to-audio environment, treating audio generation as a conditional language modeling task. Building on this, AudioLDM-2 [44] introduces a unified audio generation framework that addresses different objectives for speech, music, and sound effects, using the “language of audio” (LoA) and a GPT-2 [54] model for translation and latent diffusion [58] for generation.

AudioLM [10] employs neural codecs [78] for long-term consistency in audio generation, mapping input audio to discrete tokens and using a language modeling task. Meanwhile, DiffSound [83], using a text encoder, decoder, vector quantized variational autoencoder (VQVAE) [32], and “vocoder”, introduces a mask-based text generation strategy (MBTG) for generating text descriptions from audio labels to address the scarcity of audio-text paired data. NaturalSpeech [70] integrates a memory component in variational autoencoders [32] to achieve human-level quality in text-to-speech systems.

Significant progress in text-to-music generation has been made by notable models such as MusicLM [2], Noise2Music [29], MusicGen [18] and MeLoDy [37]. MusicLM [2] uses contrastive pretraining to align music and language embeddings, improving performance without explicit text using a semantic modeling stage based on w2v-BERT [17]. MusicGen [18] adopts a language modeling approach with melodic conditioning, enhancing control over output. MeLoDy [37] employs dual-path diffusion-based modeling for computational efficiency in music generation.

The image-to-audio and video-to-audio domains have also seen significant advancements. For instance, Im2Wav [66] generates audio from images or image sequences, using Transformers [77] to manipulate hierarchical discrete audio representation derived from VQ-VAE [76]. Similarly, based on GANs [26]: SpecVQGAN [31] enhances visually driven audio synthesis from videos, while RVQGAN [36] achieves high-fidelity audio compression. [7] introduces aligned representations that help in cross-modal retrieval and transferring between audio, visual, and text modalities. Furthermore, AV-NeRF [39] synthesizes new videos with spatial audio along different camera trajectories, and AV-RIR [56] estimates audio-visual room impulse responses from images and audio.

Multi-modal generative models of note include CoDi(-2) [9, 71]. CoDi [9] generates multiple output modalities from various input combinations using a composable generation strategy. CoDi-2 [71] extends this with enhanced modality-language alignment for encoding and generative tasks, improving zero-shot capabilities in multi-modal generation, reasoning, and interactive conversations.

Recent works like MAGNeT [93] introduce a masked generative sequence modeling method for direct audio token generation, achieving higher efficiency in text-to-music and text-to-audio tasks. DIFF-FOLEY [45] improves video-to-audio synthesis quality, relevance, and synchronization using latent diffusion and contrastive audio-visual pretraining. Affusion [82] enhances text-to-audio generation by leveraging text-to-image diffusion models, producing high-quality audio closely aligned with text prompts. These advancements significantly expand the possibilities of audio generation.

2.2 Acoustic Matching

Recent advancements in acoustic matching integrate visual and audio cues to enhance realism and immersion in generated content. AViTAR [14] performs visual-acoustic matching using a cross-modal transformer to integrate visual properties into audio and learning from in-the-wild videos. AdVerb [16] uses a geometry-aware cross-modal transformer to enhance reverberation by leveraging visual cues, significantly outperforming traditional methods. SoundSpaces 2.0 [15] provides geometry-based audio rendering for 3D environments, supporting tasks like audio-visual navigation and source localization with high fidelity and realism.

2.3 Construction of Surround Sound

In 1881, Ader used a pair of microphones and telephone receivers to relay the sound signals at the Paris Opera [60], a first-of-its-kind spatial (stereophonic) sound demonstration. In the past three decades, the development of digital signal processing [34], computer, and internet technologies have rapidly evolved spatial sound [81]. Together, they have made spatial audio commonplace now [88].

In the context of virtual environments generated by game engines, spatial audio is crucial to provide an immersive experience [12]. Traditionally, sounds in virtual worlds are created by recording individual sound sources separately and then blending them according to the intended scene setup [48] using various audio processing techniques [25, 19]. But this recording process is costly [88] and optimizing it involves generating spatial audio. While spatializing mono-audio is a popular area of research (discussed below), to the best of our knowledge, there have not been any approaches to generate spatial audio directly.

2.4 Audio Spatialization

Recent advancements in audio spatialization have focused on integrating audio and visual modalities to enhance realism. Sep-Stereo [91] uses an associative pyramid structure to generate stereo output by combining audio and visual information. Similarly, LAVSS [85] introduces a location-guided audio-visual spatial audio separator, highlighting the importance of visual cues in audio generation.

Several studies have synchronized visual cues and encoders to produce natural sound for dynamic video scenarios [92]. Methods like those in [22, 23] use a U-Net to encode monaural input and decode binaural output with visual guidance.

Other approaches, such as [21, 35] generate audio samples conditioned on diverse inputs such as text, motion key points, and position information [24, 51], respectively. Techniques like [47] utilize end-to-end trainable neural networks for sound source separation and localization in 360-degree videos. Additionally, works like [38] synthesize directional impulse responses by exploiting the spatial and directional characteristics of room reverberation.

2.5 Audio Visual Learning

In recent years, there has been a growing interest in audio-visual learning, focusing on understanding the inherent connection between visual and auditory stimuli within vision research. Noteworthy works have delved into explaining the natural synchronization between visual and audio elements [4, 6, 50]. This approach finds application in diverse domains, including speech recognition utilizing both audio and visual cues [69, 27, 86, 90], event localization through joint analysis of audio and visual information [74, 73, 80], and sound source localization [5, 28, 59, 64, 74, 89]. Additionally, there have been explorations in the use of audio-visual data for unsupervised representation learning [49, 50].

In contrast to these modalities, our focus diverges as we undertake a distinct task: the direct generation of surround sound spatial audio from images.

3 Preliminary: Conditional LDM

Latent Diffusion models (LDMs) [58] project input images into a lower-dimensional latent space. Unlike traditional diffusion models that operate in a high-dimensional pixel space, LDMs perform forward and backward diffusion processes in a lower-dimensional learned latent space. For generating

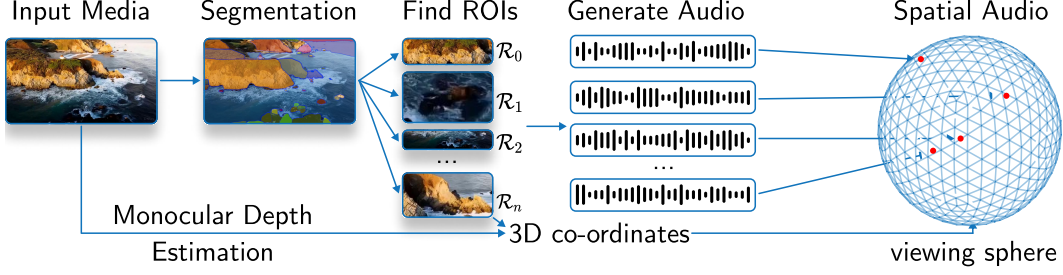


Figure 2: We show an overview of how our approach works. First, we perform image segmentation and find regions of interest for different objects in the image. We also perform monocular depth estimation on the image and produce 3D co-ordinate estimating positions for each object in the scene. We then generate mono-audio from the regions of interest. Finally, we place all of these audio according to the 3D co-ordinates on a viewing sphere producing spatial or surround sound.

audio, we have a conditional LDM ϵ_θ , which can be interpreted as a sequence of denoising autoencoders $\epsilon_\theta(\mathbf{z}_t, t, \dots)$; $t \sim \mathcal{U}(1, \dots, T)$, trained using a squared error loss (representing the denoising term in the ELBO) to denoise some variably noisy audio signals using the following objective,

$$\mathbb{E}_{\mathcal{E}(\mathbf{x}), \mathbf{c}, \epsilon \sim \mathcal{N}(0, I), t} [\|\epsilon - \epsilon_\theta(\mathbf{z}_t, t, \tau_\theta(\mathbf{c}))\|_2^2], \quad (1)$$

where \mathbf{x}_t is some noisy audio signal; \mathbf{c} is a conditioning vector; τ_θ is an encoder that projects \mathbf{c} to an intermediate representation; \mathcal{E} encodes \mathbf{x}_t into a latent representation \mathbf{z}_t ; and we jointly optimize ϵ_θ and τ_θ .

4 Methodology

Given a casually captured or generated image, animated image, or video, optionally with a text prompt that describes the content of the media, our objective is to generate spatial audio to accompany the media. We do not impose any qualitative restrictions on the way the media is captured. Our approach, SEE-2-SOUND is a compositional model that has three components (Figure 2): (1) estimating positions for regions of interest (§ 4.1); (2) generating audio for regions of interest (§ 4.2); and (3) generating spatial audio (§ 4.3).

4.1 Source Estimation

We have a visual input, for brevity let $\mathbf{I} \in \mathbb{R}^{H \times W \times C}$, represent a (H, W) -resolution and C channel image. Here, we are interested in learning from \mathbf{I} , a set of 3D coordinates, $\mathbf{x} = (x, y, z)$ representing the positions of regions of interest.

Identifying regions of interest. We first focus on finding out regions of interest from the image, \mathbf{I} . We do this by finding a segmentation mask $\mathcal{M} \in \mathbb{R}^{H \times W}$, using a pre-trained Segment Anything model [33]. From the segmentation mask \mathcal{M} , we extract the contours of each region of interest. Let $\mathcal{C} = \{\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_n\}$ represent the set of contours, where each contour \mathbf{c}_k is a sequence of points outlining one region of interest. We then compute the smallest polygon \mathcal{P}_k that bounds each contour \mathbf{c}_k . Formally, we solve the following optimization problem for each contour,

$$\min_{\mathcal{P}_k} \text{Area}(\mathcal{P}_k), \text{ subject to, } \forall \mathbf{c}_k \in \mathcal{C}, \quad (2)$$

where $\text{Area}(\mathcal{P}_k)$ denotes the area of the polygon \mathcal{P}_k , ensuring that the polygon fully encloses the contour \mathbf{c}_k . This allows us to generate a set of minimal bounding polygons, $\mathcal{P} = \{\mathcal{P}_1, \mathcal{P}_2, \dots, \mathcal{P}_k\}$.

We can also modify this to instead set each polygon to an axis-aligned bounding box (AABB) \mathcal{R}_k for each contour \mathbf{c}_k . The AABB is defined by the minimum and maximum x and y coordinates of the points in \mathbf{c}_k . Formally, for each contour \mathbf{c}_k , we define \mathcal{R}_k as,

$$\mathcal{R}_k = \left[\min(\mathbf{c}_k^{(x)}), \max(\mathbf{c}_k^{(x)}), \min(\mathbf{c}_k^{(y)}), \max(\mathbf{c}_k^{(y)}) \right]. \quad (3)$$

This allows us to generate a set of AABBs, $\mathcal{R} = \{\mathcal{R}_1, \mathcal{R}_2, \dots, \mathcal{R}_k\}$.




Figure 3: **Overview of the simulation.** We show a schematic of the virtual “ShoeBox” room utilized in our 5.1 surround sound simulation. The diagram presents the three-dimensional arrangement of multiple audio sources (labeled Source 1, Source 2, and Source 3) and the standardized microphone placement for the channels Front-Left (FL), Front-Right (FR), Center (C), Surround-Left (SL), Surround-Right (SR), and Low-Frequency Effects (LFE).

Estimate monocular depth. For the image \mathbf{I} , we also estimate the monocular depth map $\mathcal{D} \in \mathbb{R}^{H \times W}$ using a pre-trained Depth Anything model [84]. Now we can construct a set of image patches \mathbf{I}_p , for the regions of interest from either \mathcal{P} or \mathcal{R} to create a set of variably sized patches, and the 3D coordinates (x, y, z) , from the image coordinates and the monocular depth coordinates illustrated in Figure 2. We find that using AABBs usually works better for generating spatial audio.

4.2 Audio Generation

We are now interested in learning from the set of image patches \mathbf{I}_p and an optional text prompt \mathbf{c} , a set of 4D representations, $(x, y, z, \mathbf{A}^{T \times F})$ representing the positions and audio. We generate mono audio clips conditioned on each of the image patches in \mathcal{R} and the text conditioning \mathbf{c} using a pre-trained CoDi model [72]. We thus have the 4D representation $(x, y, z, \mathbf{A}^{T \times F})$ for each of the individual image patches in $\{\mathcal{R}_1, \mathcal{R}_2, \mathcal{R}_3, \dots\}$.

4.3 Generating 5.1 Surround Sound Spatial Audio

The goal of surround sound spatial audio generation is to generate ambisonic channels $\phi_N(t)$ from different non-spatial audios and positions $(x, y, z, \mathbf{A}^{T \times F})$. We first construct a virtual acoustic environment, specifically a “ShoeBox” room [62]. We place the listener centrally at the ground level of this room, serving as a reference point for audio projection and perception. For the source and microphone placements, we adhere to the ITU-R BS 775 configuration [3], which is essential for aligning our simulations with the standard 5.1 Surround Sound setups. This configuration includes designated positions for the Front-Left (FL), Front-Right (FR), Center (C), Surround-Left (SL), Surround-Right (SR), and Low-Frequency Effects (LFE) channels [65].

A critical component of our simulation involves computing the Room Impulse Responses (RIRs) using the Image Source Method (ISM) [3]. We now calculate the impulse response $h(t)$ by accounting for contributions from multiple virtual sources that arise from reflections off the room’s boundaries. For a source at position \mathbf{s} and a receiver at position \mathbf{r} , the impulse response is expressed as

$$h(t) = \sum_m \alpha_m \delta(t - \tau_m), \quad (4)$$

where $\tau_m = \frac{\|\mathbf{r} - \mathbf{s}_m\|}{c}$ denotes the time delay; c represents the speed of sound, and α_m are the attenuation factors due to reflections. This model effectively captures both the early reflections and the reverberations, which are crucial for rendering a realistic acoustic environment.

(+ AViTAR [14])	MFCC-DTW ↓	ZCR ↑	Chroma ↑	Spect ↑
CoDi [72] (baseline)	0.8×10^{-3}	0.80	0.70	0.85
SEE-2-SOUND (ours)	0.03×10^{-3}	0.95	0.77	0.95

Table 1: **Quantitative Evaluation Results.** We report average similarity scores between different audios. The similarity metrics are calculated across the method and AViTAR [14] enhanced audios of the methods. Lower MFCC-DTW scores and higher ZCR, Chroma, and Spectral scores indicate higher similarity.

After computing the RIRs, we convolve these responses with the original mono audio tracks to synthesize the spatial audio for each of the 5.1 channels. The convolution integrates the spatial effects into the audio, effectively distributing the sound across the designated channels. The synthesized signal for each channel ch is expressed as

$$\phi_{ch}(t) = \mathbf{A}^{T \times F} * h_{ch}(t), \quad (5)$$

where $\mathbf{A}^{T \times F}$ denotes the original audio signal in the time-frequency domain. Our simulation also incorporates the interference among audio sources which is a significant factor that differentiates the outcome from scenarios where fewer sources are used and then combined. Thus, the resulting sound field in a full simulation exhibits unique characteristics that cannot be replicated by simply summing the outputs of simulations with fewer sources. This interference results in spatial audio properties that are intrinsically linked to the arrangement (x, y, z) and the number of sources in the complete setup.

5 Experiments

To the best of our knowledge, no previous techniques address the problem of spatial sound generation, thus we propose a novel quantitative evaluation technique based on marginal scene guidance (Section 5.1) alongside human evaluation (Section 5.2).

5.1 Quantitative Evaluation

We conduct a quantitative evaluation of the scene guidance of generations produced by our model using various audio similarity metrics. To construct this setup, we first sample images from Laion400M [63] and the web. We then generate audio from these images using our method as well as CoDi [72] which serves as a baseline. We then use the audio outputs and images to generate modified scene-guided audio generated using AViTAR [14]. The use of an unconventional model-based evaluation stems from the lack of existing evaluations or any strong baselines. We report the average audio similarity scores between multiple such generated audio in Table 1. In particular, SEE-2-SOUND consistently generates spatial audio that matches the scene’s content as shown by AViTAR [14] only slightly altering the generated audios.

5.2 Human Evaluation

Due to the absence of any traditional metrics, we also use human evaluation to evaluate our approach. To promote future research in this field, we will open-source the evaluation protocol for human evolution. We conduct human evaluations by surveying 43 human evaluators, for which we use randomly selected images and videos from the internet that are non-iconic in nature. We use this qualitative evaluation technique specifically due to the absence of any baseline methods for spatial generation. Additionally, the task of spatial audio generation is challenging since there are multiple possible ground truth spatial audios for a given image or video. This makes it difficult for us to use any kind of similarity metric. Existing similarity metrics like LPIPS [87] or CLIP score [52] and their audio equivalents that make use of embeddings are also not feasible since we do not train embeddings directly for the spatial generation problem.

Our human evaluation task has the following key components.

<i>Metric</i>	<i>p-value (\downarrow)</i>
Spatial Audio Quality	
Realism	4.174×10^{-5}
Immersion	1.036×10^{-4}
Accuracy	1.036×10^{-4}
Clarity	5.927×10^{-6}
Consistency	2.089×10^{-6}
Audio Identification	
Overall Localization	4.523×10^{-4}
Audio Direction Identification	2.880×10^{-7}
Distance Identification	3.103×10^{-2}
Adversarial Identification	9.345×10^{-2}
Audio Matching	7.431×10^{-5}

Table 2: **Human Evaluation Results.** We evaluate our model performance across metrics identified through a user study. The hypothesis of our test was that the responses follow a uniform distribution. For the quality questions, we considered 3, 4, and 5 to be correct on a scale of 1-5. For the distance and localization problems, we used the patches and coordinates generated by our approach as correct answers. For adversarial identification and matching problems, we knew the correct answers by the nature of the problem. The significance of human preference is evaluated with χ^2 -tests at $p < 0.05$.

Spatial Audio Quality. Human evaluators qualitatively rate the realism, immersion, and accuracy of two images or videos with spatial audio generated by our method. These aspects are measured using a semantic differential scale.

Direction Identification. Evaluators estimate the sensation of direction and distance of an object given the generated spatial audio. This assesses the effectiveness of the spatial sound in conveying the sensation of sound originating from multiple directions.

Adversarial Identification. Evaluators are given a random mix of generated spatial audios and mono audios. They are tasked with discriminating whether the given audio is indeed spatial audio, their accuracy is measured.

Audio Quality. Evaluators are provided with three clips of spatial audio and three images or videos. They perform a matching task to pair the audio clips with the corresponding visual content, and their matching accuracy is measured.

More details on the tasks in our user study can be found in Appendix B. We evaluate the results for each of these individual sub-tasks and questions and conduct χ^2 -tests to evaluate significance at $p < 0.05$ level.

Furthermore, we demonstrate our outputs in the supplementary from which one can observe that SEE-2-SOUND produces compelling spatial audio that aligns well with videos and images which are either generated from video generation models or taken from the internet. We perform the user testing on 43 human evaluators and report the metrics for our human evaluation in Table 2. Users indicated a statistically significant preference towards SEE-2-SOUND in terms of spatial audio quality, directions of sound, and audio quality.

6 Discussion

Our work introduces a method for generating spatial sound from images and videos. We envision many avenues for potential impact based on our technique, including applications related to enhancing images and videos generated by learned methods, making real images interactive, human-computer interaction problems, or for accessibility.

Limitations. Although SEE-2-SOUND produces compelling spatial audio clips, there are several limitations and avenues for future work. First, our framework may not detect all of the fine details in images and videos, nor does it produce audio for every fine detail. If we try to produce audio for

many fine structures with this framework using upscaling or other zooming-in methods, we observe a lot of interference within the audios.

Second, SEE-2-SOUND does not generate audios based on any motion cues but solely based on the image priors. We could benefit from having separate appearance and motion backbones, which may potentially allow us to use any motion information to refine the generated audio.

Third, currently, our method is not able to work in real-time on an Nvidia A100-80 GPU on which the experiments were conducted. However, since SEE-2-SOUND relies on composing off-the-shelf models, we could easily swap out the models we use for other models that solve the individual sub-problems. We can easily incorporate advances in any of the subproblem areas within our framework to potentially get to real-time capabilities.

Lastly, inherent in the capabilities of generative models such as SEE-2-SOUND – which excel in producing synchronized and realistic outputs – is a significant concern regarding the potential generation and dissemination of deepfakes. The exploitation of this technology by malicious entities poses a notable risk even though we do not generate speech, as they could fabricate realistic counterfeit audio clips. We condemn any such applications and hope to be able to work on safeguards for such approaches in the future.

7 Conclusion

Our work builds an audio-generation pipeline, SEE-2-SOUND, that can generate spatial audio from images, dynamic images, or videos. We achieve this by framing it as a compositional problem and propose a novel framework to use off-the-shelf methods for this challenging new task. To the best of our knowledge, our work is also the first to tackle the problem of performing spatial audio generation. Evaluating approaches for this novel task is very challenging thus apart from doing a user study we also propose a novel quantitative evaluation method based on measuring scene guidance from audios. We and our human evaluators observe compelling outputs for images and videos both from the internet and produced by generative models. Our method is also zero-shot in nature which makes our method significantly easy to use and extensible. To this extent, we also envision users being able to swap out the models we use for advances and other models in each of the individual sub-problems in our pipeline.

We hope SEE-2-SOUND and the future works we inspire lead to the generation of truly immersive digital content. We also look forward to seeing other research on the problem we describe in this work. Our codes, models and results are open-sourced to facilitate future research in this area.

Acknowledgements

We thank Bhavya Bhatt of the University of Toronto for his help on surveying the existing literature. This research was enabled in part by support provided by the Digital Research Alliance of Canada ¹. This research was supported in part with Cloud TPUs from Google’s TPU Research Cloud (TRC) ². The resources used to prepare this research were provided, in part, by the Province of Ontario, the Government of Canada through CIFAR, and companies sponsoring the Vector Institute ³.

¹<https://alliancecan.ca/>

²<https://sites.research.google/trc/about/>

³<https://vectorinstitute.ai/partnerships/current-partners/>

References

- [1] J. Achiam, S. Adler, S. Agarwal, L. Ahmad, I. Akkaya, F. L. Aleman, D. Almeida, J. Altenschmidt, S. Altman, S. Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [2] A. Agostinelli, T. I. Denk, Z. Borsos, J. Engel, M. Verzetti, A. Caillon, Q. Huang, A. Jansen, A. Roberts, M. Tagliasacchi, et al. Musiclm: Generating music from text. *arXiv preprint arXiv:2301.11325*, 2023.
- [3] J. B. Allen and D. A. Berkley. Image method for efficiently simulating small-room acoustics. *The Journal of the Acoustical Society of America*, 65(4):943–950, 1979.
- [4] R. Arandjelovic and A. Zisserman. Look, listen and learn. In *Proceedings of the IEEE international conference on computer vision*, pages 609–617, 2017.
- [5] R. Arandjelovic and A. Zisserman. Objects that sound. In *Proceedings of the European conference on computer vision (ECCV)*, pages 435–451, 2018.
- [6] Y. Aytar, C. Vondrick, and A. Torralba. Soundnet: Learning sound representations from unlabeled video. *Advances in neural information processing systems*, 29, 2016.
- [7] Y. Aytar, C. Vondrick, and A. Torralba. See, hear, and read: Deep aligned representations. *arXiv preprint arXiv:1706.00932*, 2017.
- [8] A. Blattmann, T. Dockhorn, S. Kulal, D. Mendelevitch, M. Kilian, D. Lorenz, Y. Levi, Z. English, V. Voleti, A. Letts, V. Jampani, and R. Rombach. Stable video diffusion: Scaling latent video diffusion models to large datasets, 2023.
- [9] A. Blattmann, R. Rombach, H. Ling, T. Dockhorn, S. W. Kim, S. Fidler, and K. Kreis. Align your latents: High-resolution video synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 22563–22575, June 2023.
- [10] Z. Borsos, R. Marinier, D. Vincent, E. Kharitonov, O. Pietquin, M. Sharifi, D. Roblek, O. Teboul, D. Grangier, M. Tagliasacchi, and N. Zeghidour. Audioldm: A language modeling approach to audio generation. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 31:2523–2533, 2023.
- [11] R. Bresin, A. de Witt, S. Papetti, M. Civolani, and F. Fontana. Expressive sonification of footstep sounds. *Proceedings of ISON*, 2010:3rd, 2010.
- [12] J. Broderick, J. Duggan, and S. Redfern. The importance of spatial audio in modern games and virtual environments. In *2018 IEEE Games, Entertainment, Media Conference (GEM)*, pages 1–9. IEEE, 2018.
- [13] T. Brooks, B. Peebles, C. Homes, W. DePue, Y. Guo, L. Jing, D. Schnurr, J. Taylor, T. Luhman, E. Luhman, C. Ng, R. Wang, and A. Ramesh. Video generation models as world simulators. 2024.
- [14] C. Chen, R. Gao, P. Calamia, and K. Grauman. Visual acoustic matching. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 18858–18868, 2022.
- [15] C. Chen, C. Schissler, S. Garg, P. Kobernik, A. Clegg, P. Calamia, D. Batra, P. Robinson, and K. Grauman. Soundspaces 2.0: A simulation platform for visual-acoustic learning. *Advances in Neural Information Processing Systems*, 35:8896–8911, 2022.
- [16] S. Chowdhury, S. Ghosh, S. Dasgupta, A. Ratnarajah, U. Tyagi, and D. Manocha. Adverb: Visually guided audio dereverberation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 7884–7896, 2023.
- [17] Y.-A. Chung, Y. Zhang, W. Han, C.-C. Chiu, J. Qin, R. Pang, and Y. Wu. W2v-bert: Combining contrastive learning and masked language modeling for self-supervised speech pre-training. In *2021 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*, pages 244–250. IEEE, 2021.
- [18] J. Copet, F. Kreuk, I. Gat, T. Remez, D. Kant, G. Synnaeve, Y. Adi, and A. Défossez. Simple and controllable music generation. *Advances in Neural Information Processing Systems*, 36, 2024.
- [19] G. Dickins and R. Kennedy. Towards optimal soundfield representation. In *Audio Engineering Society Convention 106*. Audio Engineering Society, 1999.
- [20] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*, 2021.

- [21] C. Gan, D. Huang, P. Chen, J. B. Tenenbaum, and A. Torralba. Foley music: Learning to generate music from videos. In *Computer Vision—ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part XI 16*, pages 758–775. Springer, 2020.
- [22] C. Gan, H. Zhao, P. Chen, D. Cox, and A. Torralba. Self-supervised moving vehicle tracking with stereo sound. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 7053–7062, 2019.
- [23] R. Gao and K. Grauman. 2.5 d visual sound. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 324–333, 2019.
- [24] R. Garg, R. Gao, and K. Grauman. Geometry-aware multi-task learning for binaural audio generation from video. In *British Machine Vision Conference*, 2021.
- [25] M. A. Gerzon. Periphony: With-height sound reproduction. *Journal of the audio engineering society*, 21(1):2–10, 1973.
- [26] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative adversarial networks. *Commun. ACM*, 63(11):139–144, oct 2020.
- [27] D. Hu, X. Li, et al. Temporal multimodal learning in audiovisual speech recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3574–3582, 2016.
- [28] D. Hu, R. Qian, M. Jiang, X. Tan, S. Wen, E. Ding, W. Lin, and D. Dou. Discriminative sounding objects localization via self-supervised audiovisual matching. *Advances in Neural Information Processing Systems*, 33:10077–10087, 2020.
- [29] Q. Huang, D. S. Park, T. Wang, T. I. Denk, A. Ly, N. Chen, Z. Zhang, Z. Zhang, J. Yu, C. Frank, et al. Noise2music: Text-conditioned music generation with diffusion models. *arXiv preprint arXiv:2302.03917*, 2023.
- [30] R. Huang, J. Huang, D. Yang, Y. Ren, L. Liu, M. Li, Z. Ye, J. Liu, X. Yin, and Z. Zhao. Make-an-audio: Text-to-audio generation with prompt-enhanced diffusion models. In A. Krause, E. Brunskill, K. Cho, B. Engelhardt, S. Sabato, and J. Scarlett, editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 13916–13932. PMLR, 23–29 Jul 2023.
- [31] V. Iashin and E. Rahtu. Taming visually guided sound generation. In *The 32st British Machine Vision Virtual Conference*. BMVA Press, 2021.
- [32] D. P. Kingma and M. Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [33] A. Kirillov, E. Mintun, N. Ravi, H. Mao, C. Rolland, L. Gustafson, T. Xiao, S. Whitehead, A. C. Berg, W.-Y. Lo, P. Dollar, and R. Girshick. Segment anything. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4015–4026, October 2023.
- [34] O. Kirkeby, P. A. Nelson, F. Orduna-Bustamante, and H. Hamada. Local sound field reproduction using digital signal processing. *The Journal of the Acoustical Society of America*, 100(3):1584–1593, 1996.
- [35] F. Kreuk, G. Synnaeve, A. Polyak, U. Singer, A. Défossez, J. Copet, D. Parikh, Y. Taigman, and Y. Adi. Audiogen: Textually guided audio generation. In *The Eleventh International Conference on Learning Representations*, 2023.
- [36] R. Kumar, P. Seetharaman, A. Luebs, I. Kumar, and K. Kumar. High-fidelity audio compression with improved rvqgan. *Advances in Neural Information Processing Systems*, 36, 2024.
- [37] M. W. Lam, Q. Tian, T. Li, Z. Yin, S. Feng, M. Tu, Y. Ji, R. Xia, M. Ma, X. Song, et al. Efficient neural music generation. *Advances in Neural Information Processing Systems*, 36, 2024.
- [38] D. Li, T. R. Langlois, and C. Zheng. Scene-aware audio for 360 videos. *ACM Transactions on Graphics (TOG)*, 37(4):1–12, 2018.
- [39] S. Liang, C. Huang, Y. Tian, A. Kumar, and C. Xu. Av-nerf: Learning neural fields for real-world audio-visual scene synthesis. *Advances in Neural Information Processing Systems*, 36, 2024.
- [40] C.-H. Lin, J. Gao, L. Tang, T. Takikawa, X. Zeng, X. Huang, K. Kreis, S. Fidler, M.-Y. Liu, and T.-Y. Lin. Magic3d: High-resolution text-to-3d content creation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 300–309, June 2023.

- [41] Y.-B. Lin and Y.-C. F. Wang. Exploiting audio-visual consistency with partial supervision for spatial audio generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 2056–2063, 2021.
- [42] H. Ling, S. W. Kim, A. Torralba, S. Fidler, and K. Kreis. Align your gaussians: Text-to-4d with dynamic 3d gaussians and composed diffusion models. *arXiv preprint arXiv:2312.13763*, 2023.
- [43] H. Liu, Z. Chen, Y. Yuan, X. Mei, X. Liu, D. Mandic, W. Wang, and M. D. Plumbley. Audioldm: text-to-audio generation with latent diffusion models. In *Proceedings of the 40th International Conference on Machine Learning*, pages 21450–21474, 2023.
- [44] H. Liu, Y. Yuan, X. Liu, X. Mei, Q. Kong, Q. Tian, Y. Wang, W. Wang, Y. Wang, and M. D. Plumbley. Audioldm 2: Learning holistic audio generation with self-supervised pretraining. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, pages 1–13, 2024.
- [45] S. Luo, C. Yan, C. Hu, and H. Zhao. Diff-foley: Synchronized video-to-audio synthesis with latent diffusion models. *Advances in Neural Information Processing Systems*, 36, 2024.
- [46] Z. Luo, D. Chen, Y. Zhang, Y. Huang, L. Wang, Y. Shen, D. Zhao, J. Zhou, and T. Tan. Videofusion: Decomposed diffusion models for high-quality video generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10209–10218, 2023.
- [47] P. Morgado, N. Nvasconcelos, T. Langlois, and O. Wang. Self-supervised generation of spatial audio for 360 deg video. In *Advances in Neural Information Processing Systems*, volume 31, 2018.
- [48] D. Murphy and F. Neff. Spatial sound for computer games and virtual reality. In *Game sound technology and player interaction: Concepts and developments*, pages 287–312. IGI Global, 2011.
- [49] A. Owens and A. A. Efros. Audio-visual scene analysis with self-supervised multisensory features. In *Proceedings of the European conference on computer vision (ECCV)*, pages 631–648, 2018.
- [50] A. Owens, J. Wu, J. H. McDermott, W. T. Freeman, and A. Torralba. Ambient sound provides supervision for visual learning. In *Computer Vision—ECCV 2016: 14th European Conference, Amsterdam, The Netherlands, October 11–14, 2016, Proceedings, Part I 14*, pages 801–816. Springer, 2016.
- [51] K. K. Parida, S. Srivastava, and G. Sharma. Beyond mono to binaural: Generating binaural audio from mono audio with depth and cross modal attention. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 3347–3356, 2022.
- [52] D. H. Park, S. Azadi, X. Liu, T. Darrell, and A. Rohrbach. Benchmark for compositional text-to-image synthesis. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 1)*, 2021.
- [53] B. Poole, A. Jain, J. T. Barron, and B. Mildenhall. Dreamfusion: Text-to-3d using 2d diffusion. In *The Eleventh International Conference on Learning Representations*, 2023.
- [54] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsupervised multitask learners. 2019.
- [55] C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.
- [56] A. Ratnarajah, S. Ghosh, S. Kumar, P. Chiniya, and D. Manocha. Av-rir: Audio-visual room impulse response estimation. *arXiv preprint arXiv:2312.00834*, 2023.
- [57] Y. Ren, C. Hu, X. Tan, T. Qin, S. Zhao, Z. Zhao, and T.-Y. Liu. Fastspeech 2: Fast and high-quality end-to-end text to speech. In *International Conference on Learning Representations*, 2021.
- [58] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10684–10695, June 2022.
- [59] A. Rouditchenko, H. Zhao, C. Gan, J. McDermott, and A. Torralba. Self-supervised audio-visual co-segmentation. In *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 2357–2361. IEEE, 2019.
- [60] F. Rumsey. *Spatial audio*. Taylor & Francis, 2012.

- [61] C. Saharia, W. Chan, S. Saxena, L. Li, J. Whang, E. L. Denton, K. Ghasemipour, R. Gontijo Lopes, B. Karagol Ayan, T. Salimans, J. Ho, D. J. Fleet, and M. Norouzi. Photorealistic text-to-image diffusion models with deep language understanding. In S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems*, volume 35, pages 36479–36494. Curran Associates, Inc., 2022.
- [62] R. Scheibler, E. Bezzam, and I. Dokmanić. Pyroomacoustics: A python package for audio room simulation and array processing algorithms. In *2018 IEEE international conference on acoustics, speech and signal processing (ICASSP)*, pages 351–355. IEEE, 2018.
- [63] C. Schuhmann, R. Beaumont, R. Vencu, C. Gordon, R. Wightman, M. Cherti, T. Coombes, A. Katta, C. Mullis, M. Wortsman, et al. Laion-5b: An open large-scale dataset for training next generation image-text models. *Advances in Neural Information Processing Systems*, 35:25278–25294, 2022.
- [64] A. Senocak, T.-H. Oh, J. Kim, M.-H. Yang, and I. S. Kweon. Learning to localize sound source in visual scenes. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 4358–4366, 2018.
- [65] B. Series. Multichannel stereophonic sound system with and without accompanying picture. *International Telecommunication Union Radio-communication Assembly*, 2010.
- [66] R. Sheffer and Y. Adi. I hear your true colors: Image guided audio generation. In *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5. IEEE, 2023.
- [67] U. Singer, A. Polyak, T. Hayes, X. Yin, J. An, S. Zhang, Q. Hu, H. Yang, O. Ashual, O. Gafni, D. Parikh, S. Gupta, and Y. Taigman. Make-a-video: Text-to-video generation without text-video data. In *The Eleventh International Conference on Learning Representations*, 2023.
- [68] U. Singer, S. Sheynin, A. Polyak, O. Ashual, I. Makarov, F. Kokkinos, N. Goyal, A. Vedaldi, D. Parikh, J. Johnson, et al. Text-to-4d dynamic scene generation. In *Proceedings of the 40th International Conference on Machine Learning*, pages 31915–31929, 2023.
- [69] J. Son Chung, A. Senior, O. Vinyals, and A. Zisserman. Lip reading sentences in the wild. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 6447–6456, 2017.
- [70] X. Tan, J. Chen, H. Liu, J. Cong, C. Zhang, Y. Liu, X. Wang, Y. Leng, Y. Yi, L. He, et al. Naturalspeech: End-to-end text-to-speech synthesis with human-level quality. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024.
- [71] Z. Tang, Z. Yang, M. Khademi, Y. Liu, C. Zhu, and M. Bansal. Codi-2: In-context, interleaved, and interactive any-to-any generation, 2023.
- [72] Z. Tang, Z. Yang, C. Zhu, M. Zeng, and M. Bansal. Any-to-any generation via composable diffusion. *Advances in Neural Information Processing Systems*, 36, 2024.
- [73] Y. Tian, D. Li, and C. Xu. Unified multisensory perception: Weakly-supervised audio-visual video parsing. In *Computer Vision—ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part III 16*, pages 436–454. Springer, 2020.
- [74] Y. Tian, J. Shi, B. Li, Z. Duan, and C. Xu. Audio-visual event localization in unconstrained videos. In *Proceedings of the European conference on computer vision (ECCV)*, pages 247–263, 2018.
- [75] A. van den Oord, S. Dieleman, H. Zen, K. Simonyan, O. Vinyals, A. Graves, N. Kalchbrenner, A. Senior, and K. Kavukcuoglu. Wavenet: A generative model for raw audio. In *Arxiv*, 2016.
- [76] A. Van Den Oord, O. Vinyals, et al. Neural discrete representation learning. *Advances in neural information processing systems*, 30, 2017.
- [77] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. u. Kaiser, and I. Polosukhin. Attention is all you need. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.
- [78] C. Wang, A. Serrano, X. Pan, K. Wolski, B. Chen, H.-P. Seidel, C. Theobalt, K. Myszkowski, and T. Leimkühler. A neural implicit representation for the image stack: Depth, all in focus, and high dynamic range. *ACM Transactions on Graphics*, 42(6), 2023.
- [79] X. Wang, H. Yuan, S. Zhang, D. Chen, J. Wang, Y. Zhang, Y. Shen, D. Zhao, and J. Zhou. Videocomposer: Compositional video synthesis with motion controllability. *Advances in Neural Information Processing Systems*, 36, 2024.

- [80] Y. Wu, L. Zhu, Y. Yan, and Y. Yang. Dual attention matching for audio-visual event localization. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 6292–6300, 2019.
- [81] B. Xie. Spatial sound-history, principle, progress and challenge. *Chinese Journal of Electronics*, 29(3):397–416, 2020.
- [82] J. Xue, Y. Deng, Y. Gao, and Y. Li. Auffusion: Leveraging the power of diffusion and large language models for text-to-audio generation. *arXiv preprint arXiv:2401.01044*, 2024.
- [83] D. Yang, J. Yu, H. Wang, W. Wang, C. Weng, Y. Zou, and D. Yu. Diffsound: Discrete diffusion model for text-to-sound generation. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 2023.
- [84] L. Yang, B. Kang, Z. Huang, X. Xu, J. Feng, and H. Zhao. Depth anything: Unleashing the power of large-scale unlabeled data. *arXiv preprint arXiv:2401.10891*, 2024.
- [85] Y. Ye, W. Yang, and Y. Tian. Lavss: Location-guided audio-visual spatial audio separation. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 5508–5519, 2024.
- [86] J. Yu, S.-X. Zhang, J. Wu, S. Ghorbani, B. Wu, S. Kang, S. Liu, X. Liu, H. Meng, and D. Yu. Audio-visual recognition of overlapped speech for the lrs2 dataset. In *ICASSP 2020-2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 6984–6988. IEEE, 2020.
- [87] R. Zhang, P. Isola, A. A. Efros, E. Shechtman, and O. Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 586–595, 2018.
- [88] W. Zhang, P. N. Samarasinghe, H. Chen, and T. D. Abhayapala. Surround by sound: A review of spatial audio recording and reproduction. *Applied Sciences*, 7(5):532, 2017.
- [89] H. Zhao, C. Gan, A. Rouditchenko, C. Vondrick, J. McDermott, and A. Torralba. The sound of pixels. In *Proceedings of the European conference on computer vision (ECCV)*, pages 570–586, 2018.
- [90] H. Zhou, Y. Liu, Z. Liu, P. Luo, and X. Wang. Talking face generation by adversarially disentangled audio-visual representation. In *Proceedings of the AAAI conference on artificial intelligence*, volume 33, pages 9299–9306, 2019.
- [91] H. Zhou, X. Xu, D. Lin, X. Wang, and Z. Liu. Sep-stereo: Visually guided stereophonic audio generation by associating source separation. In *Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part XII 16*, pages 52–69. Springer, 2020.
- [92] Y. Zhou, Z. Wang, C. Fang, T. Bui, and T. L. Berg. Visual to sound: Generating natural sound for videos in the wild. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3550–3558, 2018.
- [93] A. Ziv, I. Gat, G. L. Lan, T. Remez, F. Kreuk, A. Défossez, J. Copet, G. Synnaeve, and Y. Adi. Masked audio generation using a single non-autoregressive transformer. *arXiv preprint arXiv:2401.04577*, 2024.

A Details of our Method

We provide the details of our approach with which we demonstrate all of our experiments and results. Our experiments ran on a single A100-80 GPU. Running our approach on an image takes $\sim 12 - 15$ seconds (non-initial runs; initial runs can take up to ~ 3 minutes) for images and ~ 5 minutes on videos that are under 20 seconds like the ones we experiment with. Our method can be easily swapped for other off-the-shelf models however we limit our experiments to these settings. Our human evaluation experiments were also performed without any text prompts to solely determine our framework’s performance.

Source Estimation. For source estimation we use a Segment Anything [33] model. We use the ViT-H/16 [20] with 14×14 windowed attention as our image encoder. Following standard approaches we rescale any inputs to 1024×1024 by rescaling and padding. For the decoder, we use the same design and weights as used by Segment Anything [33]. For monocular depth estimation, we use a Depth Anything [84] model. We use ViT-L/14 [20] as our image encoder.

Audio generation. For audio generation we use a CoDi [72] model and use the Video LDM, Audio LDM, and Text LDM components. For inference, we sample the model through 500 diffusion steps.

Generating 5.1 Surround Sound. We use a sampling rate of 16000 Hz for audio processing. For simulating the acoustic environment, we utilize a “ShoeBox” room model from Pyroomacoustics [62]. The room dimensions are dynamically calculated based on the image size. Specifically, the dimensions width $w = \text{image_width}$, height $h = \text{image_height}$, and length $l = 0.5 \cdot \text{image_width}$. In our virtual room, the microphones are placed at the following coordinates for each channel,

- Front-Left (FL): [0, 0, 100]
- Front-Right (FR): [0, 0, -100]
- Center (C): [100, 0, 0]
- Low-Frequency Effects (LFE): [0, -100, 0]
- Surround-Left (SL): [-100, 0, 0]
- Surround-Right (SR): [0, 100, 0]

B User Study Protocol

We demonstrate the user study that the participants were given. We however removed any YouTube video links we had as well as a consent form and institutional details here.

B.1 Consent Form

Your decision to participate is voluntary and you are free to withdraw at any time without any consequences. You are not required to answer any questions you do not feel comfortable answering.

Please ensure that you understand the following:

1. The purpose of the survey is to evaluate AI for spatial audio generation from image inputs and optional text prompts for the purpose of improving the quality of the outputs from these systems.
2. I will be asked a questionnaire consisting of comparisons between pairs of images, audios, and videos.
3. The questionnaire takes approximately 15 minutes to complete, and will be conducted online.
4. I may withdraw from the questionnaire at any time by exiting the questionnaire screen. If I choose to withdraw, all partial data collected will be discarded.
5. The researchers do not foresee any risks or stresses on the participant beyond what one might experience in a typical day of work.

6. The researchers also do not foresee any immediate benefits for the participant beyond what one might experience in a typical day of work.
7. The results collected will help researchers develop systems for using AI models to generate spatial audio from image inputs and optional text prompts.
8. I agree to have my responses documented and recorded digitally.
9. All data collected from the questionnaire will be kept secure. Only the researchers involved in this study will have access to the information I provide. All data collected will be destroyed 5 years after publication of results.
10. Data collected may be used in research journals, conferences, or other scholarly activities. In these cases, an anonymous subject identifier will be used and no identifying information will be provided to the audience.
11. I am at least 18 years of age.
12. I am free to ask questions about the process at any time. I can ask questions in person, or by contacting through email.
13. If requested, I will receive a copy of this form for my records.

The research study you are participating in may be reviewed for quality assurance to make sure that the required laws and guidelines are followed. If chosen, (a) representative(s) of the Human Research Ethics Program (HREP) may access study-related data and/or consent materials as part of the review. All information accessed by the HREP will be upheld to the same level of confidentiality that has been stated by the research team. If you have any questions about your rights as a participant, please contact the Ethics Review Office at (this email) or (this phone number).

B.2 Part 1: Familiarization

Before we begin, please note that the optimal testing condition involves 5.1 surround speakers. However, for this test, we ask participants to use headphones.

Watch this preselected video with correct spatial audio to familiarize yourself with the spatial audio experience.

B.3 Part 2: Rating Spatial Audio with Images

You will be exposed to images paired with corresponding spatial audio. Some of these may be generated by our system.

Rate the realism, immersion, accuracy, clarity, and consistency of the audio representation in relation to the visual content on a scale of 1 to 5, where 1 is not realistic/immersive/accurate/clear/consistent, and 5 is very realistic/immersive /accurate/clear/consistent.

Here is how we have defined these terms:

- **Realism:** How authentically the spatial audio mirrors real-world sound distribution.
- **Immersion:** The extent to which spatial audio engages and envelops the listener.
- **Accuracy:** The precision of alignment between spatial audio and visual content.
- **Clarity:** The degree to which individual sound elements are distinct and perceptible.
- **Consistency:** The uniformity of spatial audio representation across different images.

Video 1

Please rate the spatial audio in Video 1 on the following parameters, where 1 is not realistic/immersive/accurate/clear/consistent, and 5 is very realistic/immersive /accurate/clear/consistent.

Video 2

Please rate the spatial audio in Video 2 on the following parameters, where 1 is not realistic/immersive/accurate/clear/consistent, and 5 is very realistic/immersive /accurate/clear/consistent.

B.4 Part 3: Direction and Distance Identification

You will be exposed to videos paired with corresponding spatial audio. Some of these may be generated by our system.

Identify the direction and distance of sound sources within the audio landscape as depicted in the videos.

Video 3

1. Where did you perceive the sound source for the trumpet to be located in the audio clip? (You can choose multiple options)
 - Left
 - Right
 - Center
 - Behind
 - In front
 - Above
 - Below
2. How would you describe the distance of the sound source for the trumpet in the audio clip?
 - Very Close
 - Moderately Close
 - Moderately Far
 - Very Far
3. Where did you perceive the sound source for the drums to be located in the audio clip? (You can choose multiple options)
 - Left
 - Right
 - Center
 - Behind
 - In front
 - Above
 - Below
4. How would you describe the distance of the sound source for the drums in the audio clip?
 - Very Close
 - Moderately Close
 - Moderately Far
 - Very Far
5. How accurately could you localize different sound sources within the audio landscape?
 - Very Accurate
 - Moderately Accurate
 - Not Accurate

B.5 Part 4: Real vs. Fake Identification

You will be shown 2 clips, one of them has spatial audio generated by and another one has mono audio of the same clips. You will be asked to identify which of them is "real" spatial audio and "fake" spatial audio.

Video 4a

Video 4b

Which of the clips is more likely to be real spatial audio?

- Video 4a
- Video 4b

B.6 Part 5: Matching Images to Spatial Audio

You will be shown three images and corresponding spatial audio. Some of these may be generated by our system. You will be asked to match each spatial audio clip to its corresponding image.

Image 1 was shown.

Spatial Audio 1

Image 2 was shown.

Spatial Audio 2

Image 3 was shown.

Spatial Audio 3

Match each spatial audio clip to its corresponding image.

C Quantitative Evaluation Metrics

The evaluation metrics we utilized are listed below:

1. **MFCC-DTW Similarity:** Measures the similarity between the Mel-Frequency Cepstral Coefficients (MFCCs) using Dynamic Time Warping (DTW) distance.
2. **Zero-Crossing Rate (ZCR) Similarity:** Measures the similarity between the zero-crossing rates of the audio files.
3. **Rhythm Similarity:** Calculates the similarity between the rhythm patterns using the Pearson correlation coefficient.
4. **Chroma Similarity:** Measures the similarity between the chroma features.
5. **Spectral Contrast Similarity:** Compares the spectral contrast of the audio files.
6. **Perceptual Similarity:** Quantifies the perceptual similarity using the Short-Time Objective Intelligibility (STOI) metric.

We excluded the perceptual and rhythm similarity metrics for the following reasons:

- **Perceptual Similarity:** Comparing spatial audio to mono audio using perceptual similarity (such as the Short-Time Objective Intelligibility (STOI) metric) is not meaningful, as the perceptual qualities are inherently different.
- **Rhythm Similarity:** Our framework does not generate music; hence, rhythm similarity is irrelevant in our context.

D Licenses

Here we list the licensing information for code bases, pre-trained models, and data used:

- Segment Anything [33] (code and weights): Apache License 2.0, <https://github.com/facebookresearch/segment-anything/tree/main>
- Depth Anything [84] (code and weights): Apache License 2.0, <https://github.com/LiheYoung/Depth-Anything/tree/main>
- CoDi [72] (code and weights): MIT License, <https://github.com/microsoft/i-Code/tree/main/i-Code-V3>
- Visual Acoustic Matching [14] (code and weights): Creative Common CC-BY-NC Licence, <https://github.com/facebookresearch/visual-acoustic-matching>
- Laion400M [63]: Creative Common CC-BY 4.0 License, <https://laion.ai/blog/laion-5b/>
- Images from web: Creative Common CC-BY 4.0 License

Our code is made public under the Apache 2.0 License.